

```
In [18]: %matplotlib inline
import numpy as np
import scipy as sp
import scipy.stats as stats
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns; sns.set()
import patsy
import sklearn
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
from sklearn.discriminant_analysis import QuadraticDiscriminantAnalysis

# Additional libs
from scipy.stats import norm
from IPython.display import display, Markdown, Latex
import scipy.stats
import sklearn.linear_model
import sklearn.discriminant_analysis
import sklearn.preprocessing
import sklearn.model_selection
import sklearn.neighbors
from sklearn import svm, metrics
import sklearn.tree
from sklearn.metrics import accuracy_score
from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix
from sklearn.ensemble import RandomForestClassifier

# from sklearn.discriminant_analysi
```

Name: Chi Huynh

Kaggle Username: cthuynh: <https://www.kaggle.com/cthuynh> (<https://www.kaggle.com/cthuynh>)

<https://www.kaggle.com/c/otto-group-product-classification-challenge/submissions> (<https://www.kaggle.com/c/otto-group-product-classification-challenge/submissions>)

Project: Otto Group Product Classification Challenge

Github: <https://github.com/chithihuynh/ottoFinalProject> (<https://github.com/chithihuynh/ottoFinalProject>)

## Data Cleaning

```
In [19]: train = pd.read_csv('train.csv')
test = pd.read_csv('test.csv')
```

- Original Data columns: I will drop the "id" column since it is just counting each data which is already done by pd.

```
In [21]: print("train columns:",train.columns)
```

```
train columns: Index(['id', 'feat_1', 'feat_2', 'feat_3', 'feat_4',  
'feat_5', 'feat_6',  
    'feat_7', 'feat_8', 'feat_9', 'feat_10', 'feat_11', 'feat_12',  
'feat_13', 'feat_14', 'feat_15', 'feat_16', 'feat_17', 'feat_18',  
'feat_19', 'feat_20', 'feat_21', 'feat_22', 'feat_23', 'feat_24',  
'feat_25', 'feat_26', 'feat_27', 'feat_28', 'feat_29', 'feat_30',  
'feat_31', 'feat_32', 'feat_33', 'feat_34', 'feat_35', 'feat_36',  
'feat_37', 'feat_38', 'feat_39', 'feat_40', 'feat_41', 'feat_42',  
'feat_43', 'feat_44', 'feat_45', 'feat_46', 'feat_47', 'feat_48',  
'feat_49', 'feat_50', 'feat_51', 'feat_52', 'feat_53', 'feat_54',  
'feat_55', 'feat_56', 'feat_57', 'feat_58', 'feat_59', 'feat_60',  
'feat_61', 'feat_62', 'feat_63', 'feat_64', 'feat_65', 'feat_66',  
'feat_67', 'feat_68', 'feat_69', 'feat_70', 'feat_71', 'feat_72',  
'feat_73', 'feat_74', 'feat_75', 'feat_76', 'feat_77', 'feat_78',  
'feat_79', 'feat_80', 'feat_81', 'feat_82', 'feat_83', 'feat_84',  
'feat_85', 'feat_86', 'feat_87', 'feat_88', 'feat_89', 'feat_90',  
    'feat_91', 'feat_92', 'feat_93', 'target'],  
    dtype='object')
```

```
In [22]: print("test columns:",test.columns)
```

```
test columns: Index(['id', 'feat_1', 'feat_2', 'feat_3', 'feat_4', 'feat_5', 'feat_6', 'feat_7', 'feat_8', 'feat_9', 'feat_10', 'feat_11', 'feat_12', 'feat_13', 'feat_14', 'feat_15', 'feat_16', 'feat_17', 'feat_18', 'feat_19', 'feat_20', 'feat_21', 'feat_22', 'feat_23', 'feat_24', 'feat_25', 'feat_26', 'feat_27', 'feat_28', 'feat_29', 'feat_30', 'feat_31', 'feat_32', 'feat_33', 'feat_34', 'feat_35', 'feat_36', 'feat_37', 'feat_38', 'feat_39', 'feat_40', 'feat_41', 'feat_42', 'feat_43', 'feat_44', 'feat_45', 'feat_46', 'feat_47', 'feat_48', 'feat_49', 'feat_50', 'feat_51', 'feat_52', 'feat_53', 'feat_54', 'feat_55', 'feat_56', 'feat_57', 'feat_58', 'feat_59', 'feat_60', 'feat_61', 'feat_62', 'feat_63', 'feat_64', 'feat_65', 'feat_66', 'feat_67', 'feat_68', 'feat_69', 'feat_70', 'feat_71', 'feat_72', 'feat_73', 'feat_74', 'feat_75', 'feat_76', 'feat_77', 'feat_78', 'feat_79', 'feat_80', 'feat_81', 'feat_82', 'feat_83', 'feat_84', 'feat_85', 'feat_86', 'feat_87', 'feat_88', 'feat_89', 'feat_90', 'feat_91', 'feat_92', 'feat_93'], dtype='object')
```

```
In [23]: # DROP: id - no need
train = train.drop(['id'], axis=1)
test = test.drop(['id'], axis=1)
```

- Change column types for 'target' column of the training data. This is the outcome column which was entered as a 'category' dtype. Therefore, the target column will be change into a 'cat code' which is an 'int 8' type for evaluation. Features will be changed from a int64 to float64 to allow decimals during the evaluation of the data.

```
In [28]: train.dtypes
```

```
Out[28]: feat_1      int64
         feat_2      int64
         feat_3      int64
         feat_4      int64
         feat_5      int64
         ...
         feat_90     int64
         feat_91     int64
         feat_92     int64
         feat_93     int64
         target      category
         Length: 94, dtype: object
```

```
In [42]: # Change to float and cat for eval
         train.target = pd.Categorical(train.target)
         yTrain = train.target.cat.codes

         xTrain = train.values[:, :-1].astype('float')
         xTest = test.values.astype('float')
```

```
In [44]: print("yTrain type:", yTrain.dtypes)
         print("xTrain (each cell) type:", (type( xTrain[0,0])))
```


```
yTrain type: int8
xTrain (each cell) type: <class 'numpy.float64'>
```

## Exploratory Data Analysis (EDA)

### Histogram of Target in Training Data

Let's look at a histogram of the target data in the training set to see how the outcomes turns out. There seems to be alot of variation between each class.

```
In [46]: plt.figure(figsize=(12,5))
sns.countplot(train['target']);
# sns.
```



A look at the KDE over the histogram to see the differences also.

```
In [27]: # Pairplots
# sns.pairplot(data=train, hue='target',)
sns.distplot(yTrain)
plt.xlabel('Classes')
plt.ylabel('Probability')
plt.title("Training Data Outcomes");
```



```
In [45]: # # TODO REMOVE?
# plt.scatter(xTrain[:,0], xTrain[:,1], c=yTrain, edgecolors='k')
# plt.show()
# plt.scatter(xTest[:,0], xTest[:,1], edgecolors='k')
# plt.show()
```

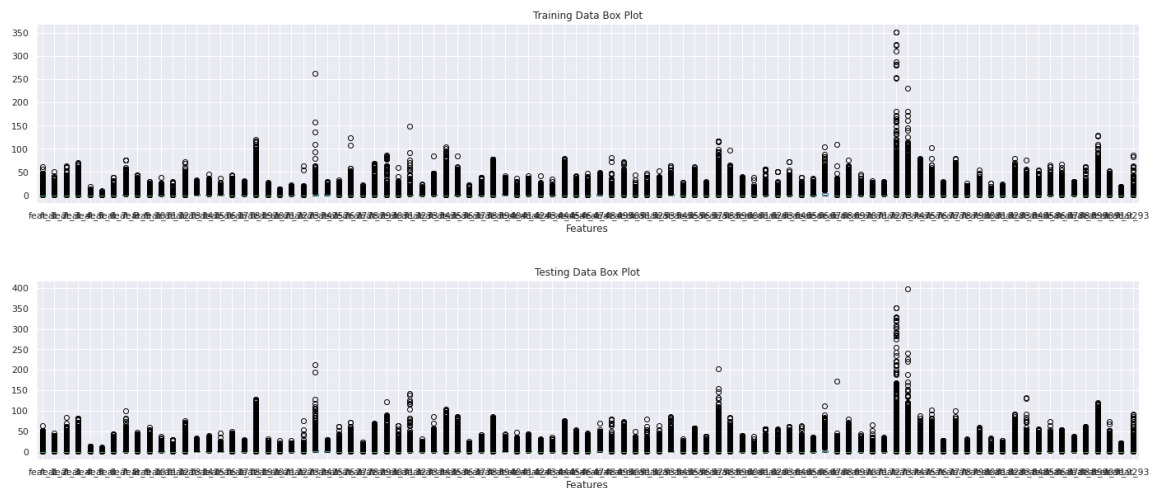
## Scaling

- From the histograms, I decided that scaling is needed since outliers can occur.
- Especially needed for KNN data since KNN uses distance metric to determine the "nearest neigh", the standard scaling will not be robust due to the outliers. There are alot of outliers:

```
In [52]: print('Train is', xTrain.shape, 'test is', xTest.shape)
```

Train is (61878, 93) test is (144368, 93)

```
In [55]: # Outliers need scaling for KNN
train.plot.box(figsize=(24,4))
plt.title('Training Data Box Plot')
plt.xlabel('Features')
plt.show()
test.plot.box(figsize=(24,4))
plt.title('Testing Data Box Plot')
plt.xlabel('Features')
plt.show()
```




```
In [50]: # Scaling train data
def scaleData(data, verbose=False):
    if verbose:
        display("Scaled From:", data, data.shape)
    data = sklearn.preprocessing.scale(data, axis=0)
    if verbose:
        display("Scaled To:", data, data.shape)
    return data
```

### Why is scaling needed?


- When scaling the data, we still keep the distribution of the values but transform it in such a way that the mean is 0 with an ideal variance of 1.

```
In [68]: # Scaling the data still keeps the distribu bit the values has  
# transform such that the mean is 0 and ideally the variance is 1  
scale_Xtrain = scaleData(xTrain)  
  
# Mean  
fig, ax = plt.subplots(nrows=1, ncols=2, figsize=(12,4))  
ax[0].plot(np.mean(xTrain, axis=0))  
plt.title('No Scaling')  
ax[1].plot(np.mean(scale_Xtrain, axis=0));  
plt.title('Mean of 0 after scaling')
```


Out[68]: Text(0.5, 1.0, 'Mean of 0 after scaling')



```
In [70]: # var num to 1  
fig, ax = plt.subplots(nrows=1, ncols=2, figsize=(12,4))  
ax[0].plot(np.var(xTrain, axis=0))  
ax[1].plot(np.var(scale_Xtrain, axis=0))  
plt.title('Variance of 1 after scaling');
```



```
In [71]: col=10
fig,ax = plt.subplots(nrows=1, ncols=2, figsize=(12,4))
sns.kdeplot(xTrain[:,col], ax=ax[0])
ax[1].set_title('Original Data');
sns.kdeplot(scale_Xtrain[:,col], ax=ax[1])
ax[1].set_title('Scaled Data');
```



## Model

Several classification methods will be use to train the data. The classification with the highest accuracy score will be use to run on the given test data set provided by Otto (test.csv).

- The training set will split into 0.3 part testing and the rest is the training set.
- KDE with a histogram will show the outcome of each classification.
- Confusion matrix for each will also be shown.

```
In [47]: # Helper function to get max predicted prob class
def getPredictProbaDistrib(predicted_proba):
    ys = []
    for i in range(len(predicted_proba)):
        y = np.argmax(predicted_proba[i])
        ys.append(y)
    return ys
```



```

In [110]: def runClassifiers(models,modelNames, dataset, test_size, test=[]):
    # Get data set and scale it
    X, y = dataset
    X = scaleData(X)[:,:]

    # Split of training and testing data
    xTrain, xTest, yTrain, yTest = sklearn.model_selection.train_test_split(X, y, random_state=0, test_size=test_size)
    if len(test) != 0:
        test = scaleData(test)[:,:]
        xTrain = X
        xTest = test

        yTrain = y
        yTest = []
    accuracy_scores = []
    cms = []
    result = []

    # for each model
    for i in range(len(models)):
        print("\n-----")
        print(modelNames[i])
        print("-----")

        # Fit and Predict
        model = %time models[i].fit(xTrain, yTrain)
        yhat = %time model.predict(xTest)
        predict_proba = %time model.predict_proba(xTest)
        yProb = getPredictProbaDistrib(predict_proba)

        # Analysis
        if len(test) == 0:
            # Analysis
            score = accuracy_score(yTest, yhat);
            print('Accuracy Score:', np.round(score, 3));
            accuracy_scores.append(score);
            cm = confusion_matrix(yTest, yhat);
            print(cm)
            cms.append(cm);
            rp = classification_report(yTest, yhat)
            print(rp)

            result.append(score)
        else:
            result.append(predict_proba)

        # Graph it
        sns.kdeplot(yTest, shade=True, color='g', Label = 'True y')
        sns.kdeplot(yhat, shade=True, color='r', Label = 'y hat')
        sns.kdeplot(yProb, shade=True, color='b', Label = 'y proba
(max)')

plt.title(str(modelNames[i]))

```

```
plt.xlabel('Classes')  
plt.ylabel('Probability')  
plt.legend(loc='best')  
plt.show();  
  
return result
```

```

In [111]: # This was separted since KNN's model needed an extra 'k' var
def runKNNClassifier(dataset, test_size, test=[]):
    # scale data
    X, y = dataset
    X = scaleData(X)[:,:]

    # get
    xTrain, xTest, yTrain, yTest = sklearn.model_selection.train_test_split(X, y, random_state=0, test_size=test_size)

    ks = [1]
    for k in ks:

        # fit and predict
        model = sklearn.neighbors.KNeighborsClassifier(n_neighbors=k, n_jobs=-1)
        model_mod = %time model.fit(xTrain,yTrain)
        yhat = %time model_mod.predict(xTest)
        predict_proba = %time model_mod.predict_proba(xTest)

        # Analysis
        score = accuracy_score(yTest, yhat);
        print('Accuracy Score:', np.round(score, 3));
        cm = confusion_matrix(yTest, yhat);
        print(cm)

        rp = classification_report(yTest, yhat)
        print(rp)
        yProb = getPredictProbaDistrib(predict_proba)

        # Graph
        plt.hist(yhat)
        plt.show()

        plt.hist(predict_proba)
        plt.show()

        sns.kdeplot(yTest, shade=True, color='g', Label = 'True y')
        sns.kdeplot(yhat, shade=True, color='r', Label = 'y hat')
        sns.kdeplot(yProb, shade=True, color='b', Label = 'y proba
(max)')

    return score, cm, yhat, predict_proba

print("----- KNN Training Model -----")
print("-----")
dataset = (xTrain, yTrain)
test_size = 0.3
knnResult = runKNNClassifier(dataset, test_size)

```

```
----- KNN Training Model -----
```

```
CPU times: user 5.12 s, sys: 12.1 ms, total: 5.13 s
```

Wall time: 5.1 s

CPU times: user 2min 31s, sys: 148 ms, total: 2min 31s

Wall time: 40.9 s

CPU times: user 2min 37s, sys: 185 ms, total: 2min 38s


Wall time: 40.9 s


Accuracy Score: 0.746

[	279	25	7	3	2	40	32	75	103]
[	6	3484	993	224	17	13	55	23	13]
[	4	948	1161	146	1	13	46	10	4]
[	3	271	188	329	5	23	12	2	3]
[	3	17	5	3	772	1	3	2	2]
[	63	46	12	34	1	3877	68	74	61]
[	34	96	96	25	8	62	504	42	12]
[	58	37	18	7	4	122	50	2201	50]
[	95	30	7	8	15	64	14	62	1236]]
			precision			recall		f1-score	support

0	0.51	0.49	0.50	566
1	0.70	0.72	0.71	4828
2	0.47	0.50	0.48	2333
3	0.42	0.39	0.41	836
4	0.94	0.96	0.95	808
5	0.92	0.92	0.92	4236
6	0.64	0.57	0.61	879
7	0.88	0.86	0.87	2547
8	0.83	0.81	0.82	1531

accuracy			0.75	18564
macro avg	0.70	0.69	0.70	18564
weighted avg	0.75	0.75	0.75	18564





Note: Only the KNN will have all three graphs. The rest will just have the last KDE graph since it give us the most information.

```
In [112]: # https://scikit-learn.org/stable/modules/tnbree.html
# tre = sklearn.tree.DecisionTreeClassifier().fit(xTrain, yTrain)

print("----- Training Models -----")
print("-----")
models = [sklearn.linear_model.LogisticRegression(multi_class='mult
inomial', penalty='l1', solver='saga', tol=0.1,max_iter=len(yTrain)), LinearDiscriminantAnalysis(), QuadraticDiscriminantAnalysis(),
sklearn.tree.DecisionTreeClassifier(), RandomForestClassifier()]
modelNameNames = ['Logistic Regression', 'Linear Discriminant Analysis', 'Quadratic Discriminant Analysis', 'Tree', 'Random Forest']
dataset = (xTrain, yTrain)
test_size = 0.3

# get best accuracy score
results = runClassifiers(models, modelNameNames, dataset, test_size);
```

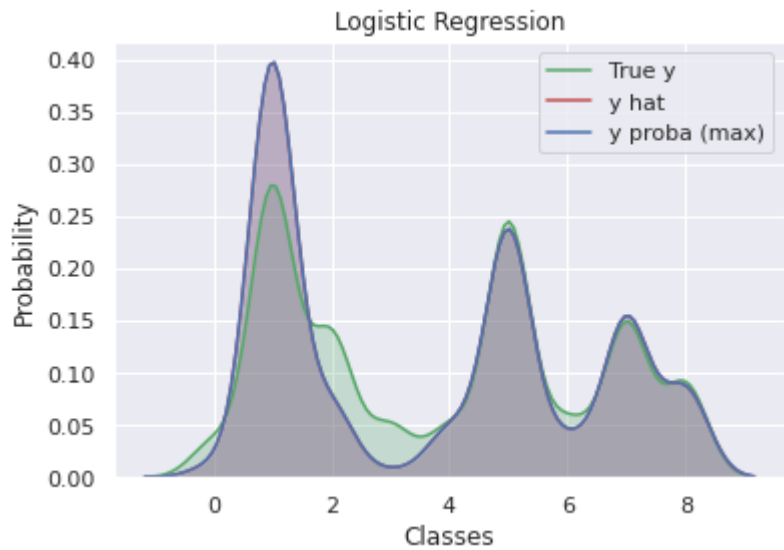
----- Training Models -----  
-----

-----  
Logistic Regression  
-----

CPU times: user 1.03 s, sys: 1  $\mu$ s, total: 1.03 s  
Wall time: 1.03 s  
CPU times: user 5.18 ms, sys: 5  $\mu$ s, total: 5.19 ms  
Wall time: 2.95 ms  
CPU times: user 10.9 ms, sys: 12  $\mu$ s, total: 10.9 ms  
Wall time: 5.32 ms  
Accuracy Score: 0.739

[[ 124	75	0	1	1	67	9	141	148]
[ 2	4316	413	4	36	12	22	16	7]
[ 0	1661	586	2	21	3	48	9	3]
[ 0	598	97	88	5	35	9	4	0]
[ 0	76	1	0	728	1	0	2	0]
[ 11	90	7	4	0	3914	52	96	62]
[ 12	202	66	4	2	83	423	81	6]
[ 24	74	10	0	3	79	18	2302	37]
[ 19	100	0	2	3	74	10	89	1234]]

	precision	recall	f1-score	support
0	0.65	0.22	0.33	566
1	0.60	0.89	0.72	4828
2	0.50	0.25	0.33	2333
3	0.84	0.11	0.19	836
4	0.91	0.90	0.91	808
5	0.92	0.92	0.92	4236
6	0.72	0.48	0.58	879
7	0.84	0.90	0.87	2547
8	0.82	0.81	0.82	1531
accuracy			0.74	18564
macro avg	0.75	0.61	0.63	18564
weighted avg	0.74	0.74	0.71	18564



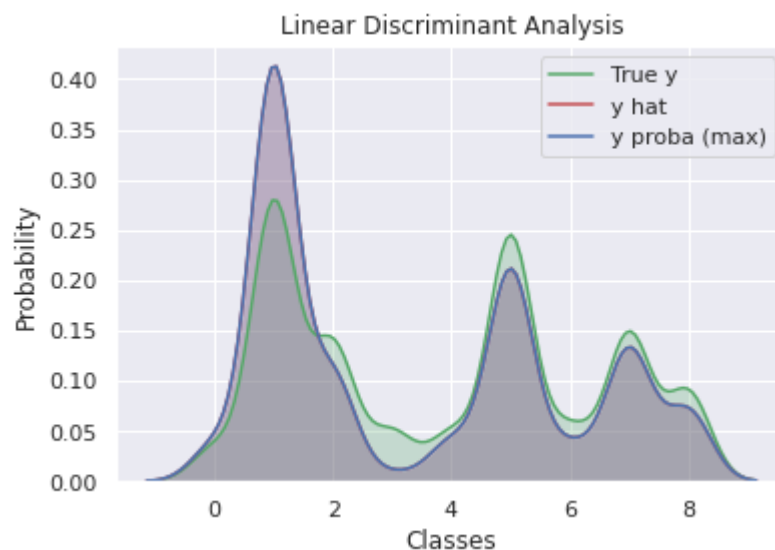
```

-----
Linear Discriminant Analysis
-----
CPU times: user 477 ms, sys: 0 ns, total: 477 ms
Wall time: 383 ms
CPU times: user 5.98 ms, sys: 0 ns, total: 5.98 ms
Wall time: 2.83 ms
CPU times: user 11 ms, sys: 4 µs, total: 11 ms
Wall time: 5.42 ms
Accuracy Score: 0.7
[[ 279  131    0    0    0   10    5   65   76]
 [   13 3954  782   19   28    2   22    7    1]
 [    2 1417  834   16   13    1   45    4    1]
 [    1  606  106   91    4   11   13    3    1]
 [    0  152    5    0  649    0    0    2    0]
 [   59  322   15    8    0 3628   57   90   57]
 [   39  281   63    4    0   29  419   43    1]
 [  105  265   22    0    1   42   16 2064   32]
 [  135  230    0    0    0   28    8   60 1070]]
      precision      recall  f1-score   support

0         0.44         0.49         0.47         566
1         0.54         0.82         0.65        4828
2         0.46         0.36         0.40        2333
3         0.66         0.11         0.19         836
4         0.93         0.80         0.86         808
5         0.97         0.86         0.91        4236
6         0.72         0.48         0.57         879
7         0.88         0.81         0.85        2547
8         0.86         0.70         0.77        1531

accuracy          0.70        18564
macro avg         0.72         0.60         0.63        18564
weighted avg      0.73         0.70         0.69        18564

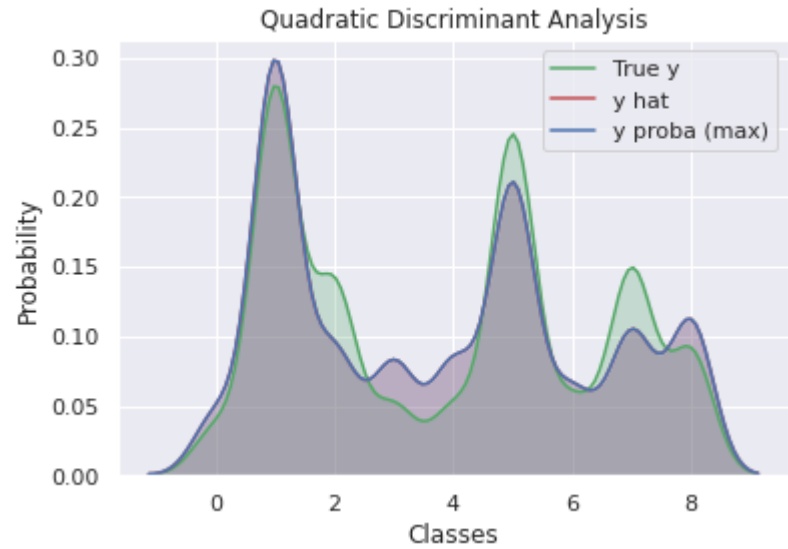
```





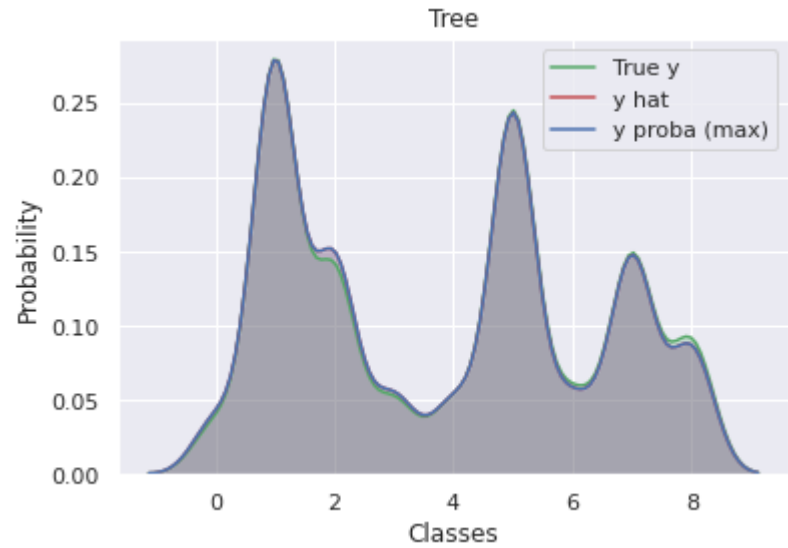
-----  
Quadratic Discriminant Analysis  
-----  
CPU times: user 252 ms, sys: 19  $\mu$ s, total: 252 ms  
Wall time: 129 ms  
CPU times: user 188 ms, sys: 0 ns, total: 188 ms  
Wall time: 93.6 ms  
CPU times: user 183 ms, sys: 4 ms, total: 187 ms  
Wall time: 93.4 ms  
Accuracy Score: 0.661

[[ 207 37 8 23 34 16 29 23 189]				
[ 10 3428 559 391 291 13 100 11 25]				
[ 1 1161 708 243 108 1 96 8 7]				
[ 0 285 41 418 54 4 23 3 8]				
[ 0 41 0 7 752 3 2 0 3]				
[ 114 73 42 165 27 3445 121 75 174]				
[ 23 108 87 91 27 39 472 21 11]				
[ 323 26 43 11 24 86 167 1598 269]				
[ 61 62 12 30 52 38 10 30 1236]]				
	precision	recall	f1-score	support
0	0.28	0.37	0.32	566
1	0.66	0.71	0.68	4828
2	0.47	0.30	0.37	2333
3	0.30	0.50	0.38	836
4	0.55	0.93	0.69	808
5	0.95	0.81	0.87	4236
6	0.46	0.54	0.50	879
7	0.90	0.63	0.74	2547
8	0.64	0.81	0.72	1531
accuracy			0.66	18564
macro avg	0.58	0.62	0.58	18564




```
-----
Tree
-----
CPU times: user 1.05 s, sys: 2 µs, total: 1.05 s
Wall time: 1.04 s
CPU times: user 6.24 ms, sys: 0 ns, total: 6.24 ms
Wall time: 5.92 ms
CPU times: user 6.19 ms, sys: 4 µs, total: 6.19 ms
Wall time: 5.98 ms
Accuracy Score: 0.708
[[ 239  26  12   3   3  48  38  88 109]
 [  24 3336 967 247  18  45  98  50  43]
 [   9  903 1133 145   2  24  75  17  25]
 [   5  236  163 366   2  25  16   8  15]
 [   4  13   8   5 766   5   2   1   4]
 [  59  77  33  54   3 3750  76 112  72]
 [  53 110 100  33   7  94 401  59  22]
 [  97  41  33   9   8 130  90 2070  69]
 [ 131  45  18  18   6  72  42 109 1090]]
      precision      recall  f1-score   support

0          0.38         0.42         0.40         566
1          0.70         0.69         0.69        4828
2          0.46         0.49         0.47        2333
3          0.42         0.44         0.43         836
4          0.94         0.95         0.94         808
5          0.89         0.89         0.89        4236
6          0.48         0.46         0.47         879
7          0.82         0.81         0.82        2547
8          0.75         0.71         0.73        1531
```



```
-----
Random Forest
-----
CPU times: user 7.39 s, sys: 102 µs, total: 7.39 s
Wall time: 7.39 s
CPU times: user 387 ms, sys: 0 ns, total: 387 ms
Wall time: 386 ms
CPU times: user 389 ms, sys: 0 ns, total: 389 ms
Wall time: 389 ms
Accuracy Score: 0.8
[[ 213   21    3    1    2   48   14  131  133]
 [   0 4257  496   24    6    9   23   11    2]
 [   0 1183 1077   30    0    3   24   12    4]
 [   0  326  117  350    5   24   10    2    2]
 [   2   20    1    0  783    2    0    0    0]
 [   9   59    3    4    0 3987   50   72   52]
 [  11  137   60   16    4   78  485   74   14]
 [  17   29    8    0    4   74   10 2377   28]
 [  25   35    1    0    3   63   11   71 1322]]
      precision      recall  f1-score   support

0          0.77        0.38        0.51        566
1          0.70        0.88        0.78       4828
2          0.61        0.46        0.53       2333
3          0.82        0.42        0.56        836
4          0.97        0.97        0.97        808
5          0.93        0.94        0.94       4236
```



```
In [119]: # Change data into form that is required to submit for kaggle
y = []
for i in range(11):
    y.append(0)

testResults = []
i = 1
submits = {}

for result in results[0]:
    # print(result)
    index = np.argmax(result)
    yi = y.copy()
    yi[0] = i
    yi[index] = 1

    testResults.append(yi)
    submits[i] = yi

    i = i + 1
```

## Results and Analysis:

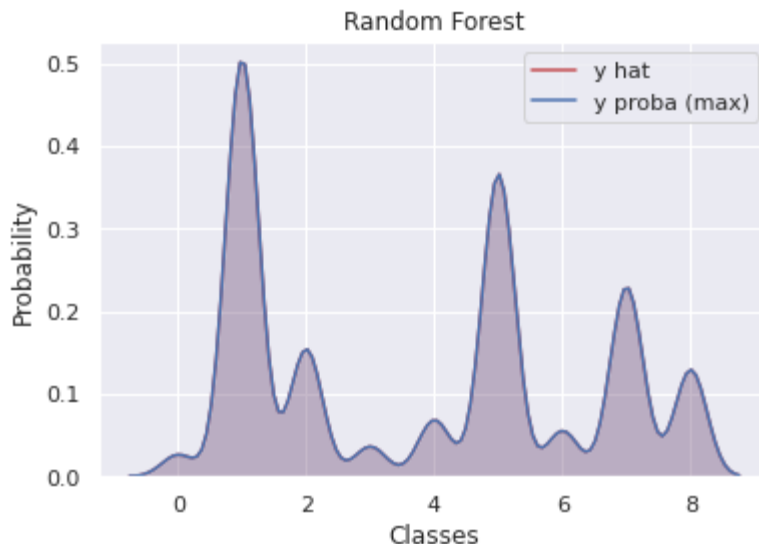
- Best result (Please see the 'Models' section for each classification results):

```
In [109]: # Run model with best accuracy score
print("----- Best Test Model -----")
print("-----")
# scores = results
bestModelIndex = np.argmax(results)
best = runClassifiers([models[bestModelIndex]], [modelName[bestModelIndex]], dataset, test_size, xTest)
```

```
----- Best Test Model -----
-----
```

```
-----
Random Forest
-----
```

```
CPU times: user 11.3 s, sys: 131 µs, total: 11.3 s
Wall time: 11.3 s
CPU times: user 3.24 s, sys: 50 µs, total: 3.24 s
Wall time: 3.24 s
CPU times: user 3.25 s, sys: 29 µs, total: 3.25 s
Wall time: 3.25 s
```



```
In [108]: print('Best model is:', modelName[bestModelIndex])
print('Score of:', results[bestModelIndex])
```

```
Best model is: Random Forest
Score of: 0.8037060978237449
```

All in all, the highest accuracy score was from the Random Forest classification. But of the ones that we mentioned in class, the KNN with 1 nearest neighbor came in second place.

The F1 score was also used to determine the best model because it mirrors the accuracy score very well. As you can see from the below model reports, for each F1 score, the Random Forest is usually higher than the other models with KNN being the second to follow it. This seems plausible since F1 score is the harmonic mean of precision and recall. Since we are looking for precise and true outcome. That is why the two evaluations were used to measure which model should be used.

## Discussion and Conclusion: 10 pts

**Background** This notebook contains the data analysis for 'Otto Group Product Classification Challenge - Classify products into the correct category'. The goal is to classify the test data set into 9 different major shopping department categories. But the categories are not known to us, it is simply stated as 'Class 1-9'.

Since Otto is a worldwide company, there has been different classification for the same item. In order to refine this classification process between the different culture's interpretations, it is vital to correctly classify the right classes.

Therefore, there are 93 features that a product can belong to. Of those 93 features, each as a number associated with how many times that feature was associated with that event (class).

**Discussion: Result** While training the different models, the result was the first and second best model was with the Random Forest then the KNN model. I was surprised on how well the overlap of the tree model looked in the graph but it did not produce the highest accuracy score. Since the KNN worked so well for  $k = 1$  nearest neighbor, I would suggest to do  $k=2$  near neighbor also. It just took a rather long time for even  $k = 1$  so I did not do it this time.

Therefore, Random Forest was used to submit to Kaggle to determine the category of each of the test data set that was provided.

In [ ]:

## Write-up: 5 pts

- Is the writeup organized and clear? (2 pts)
- Are the codes commented and organized? (2 pts)
- Did author use git? (1 pts)

In [ ]: